



# Akka Streams, Kafka, Kinesis

# Peter Vandenabeele

## Mechelen, June 25, 2015

# #StreamProcessingBe



# **whoami : Peter Vandenabeele**

@peter\_v

# @All\_Things\_Data (my consultancy)

# current client:

# Real Impact Analytics @RIAnalytics

# Telecom Analytics (emerging markets)

# Agenda



## 5' Intro (Peter)

# 40' Akka Streams, Kafka, Kinesis (Peter)

# 45' Spark Streaming and Kafka Demo (Gerard)

## 15' Open discussion (all)

# 30' beers (doors close at 21:30)



# Many thanks to

# @All\_Things\_Data (beer)

@maasg (Gerard Maas)

you !

=> Note: always looking for locations



# Agenda

# Why ?

# Akka + Akka Streams

## Demo

# Kafka + Kinesis

## Demo

# Why !



# Why ?

# distributed state

# distributed failure

# slow consumers



# Akka

# Why (for iHeartRadio



# Why we picked Akka cluster as the core architecture for our microservice

- Out-of-the-box clustering infrastructure
  - Loose coupling without the cost of JSON parsing.
  - Transparent programming model within and across microservices.
  - Strong community and commercial support
  - High resiliency, performance and scalability.

source: [tech.iheart.com/post/121599571574/why-we-picked-akka-cluster-as-our-microservice](http://tech.iheart.com/post/121599571574/why-we-picked-akka-cluster-as-our-microservice)

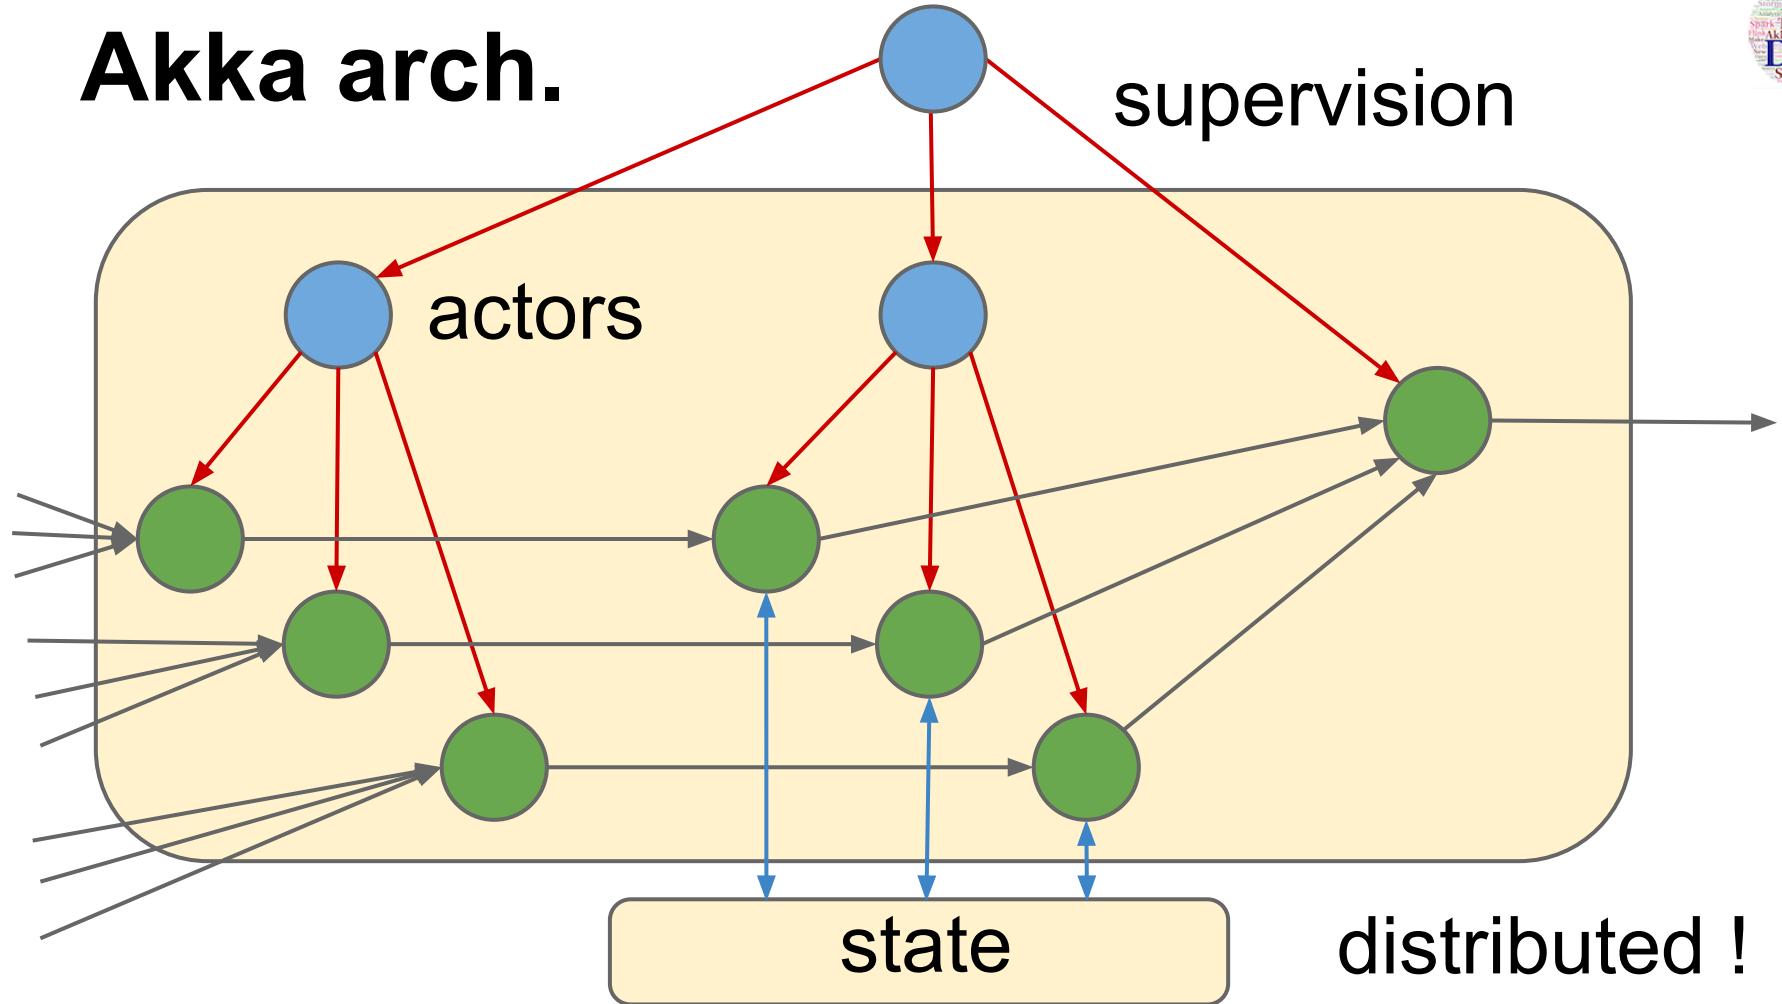
# Akka design



# Building

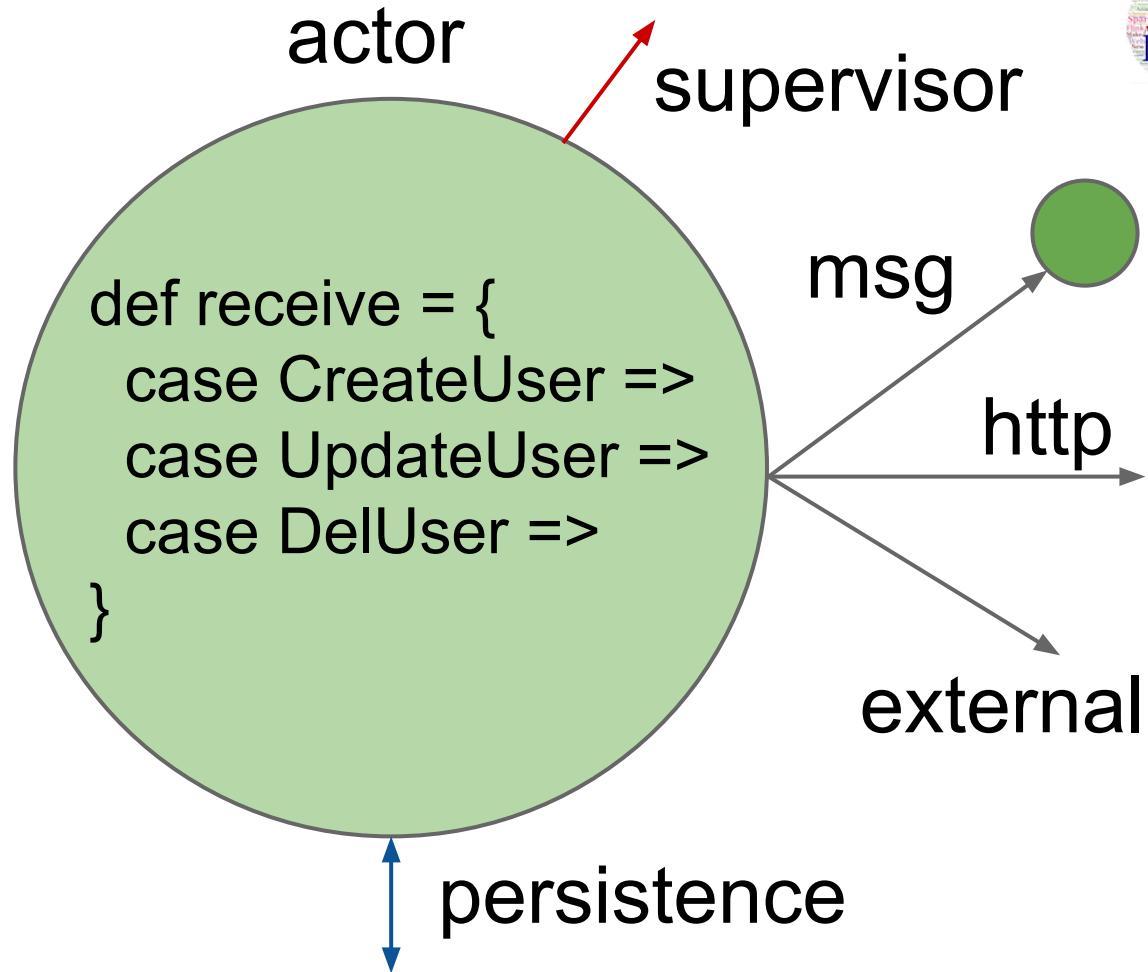
- concurrent       $\leq$  many (slow) CPU's
  - distributed       $\leq$  distributed state
  - resilient         $\leq$  distributed failure
  - applications     $\leq$  platform
  - on JVM             $\leq$  Erlang OTP

# Akka arch.



# Akka actor

- msg
- | 4 | 3 | 2 | 1 |
- msgs are sent
- recv'd in order
- single thread
- stateful !
- errors go “up”





# Akka usage + courses

- concurrent programming *not* easy ...
- but *without* Akka ... would be much harder
- Spark (see log extract next slide)
- Flink (version 0.9 of 24 June)
- local projects (e.g."Wegen en verkeer")
- BeScala Meetup now runs Akka intro course
- commercial courses (Cronos, Scala World...)



# Spark heavily based on Akka

## log extract from Spark:

**java -cp ...**

```
"org.apache.spark.executor.CoarseGrainedExecutorBackend"
"--driver-url" "akka.tcp://sparkDriver@docker-master:
51810/user/CoarseGrainedScheduler"
"--executor-id" "23" "--hostname" "docker-slave2" "--cores" "8"
"--worker-url" "akka.tcp://sparkWorker@docker-slave2:
50268/user/Worker"
```



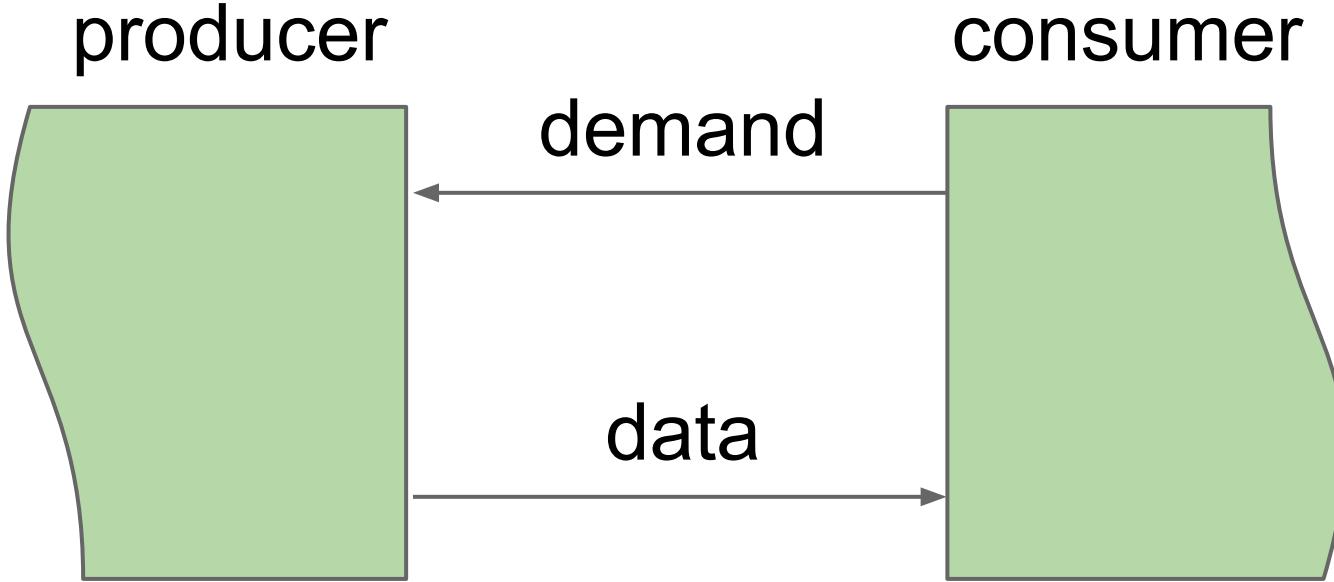
# Akka Streams



# Reactive Streams

- <http://reactive-streams.org>
  - exchange of  
*stream data* across  
asynchronous boundary in  
*bounded fashion*
  - building and industry standard (open IP)

# Demand based



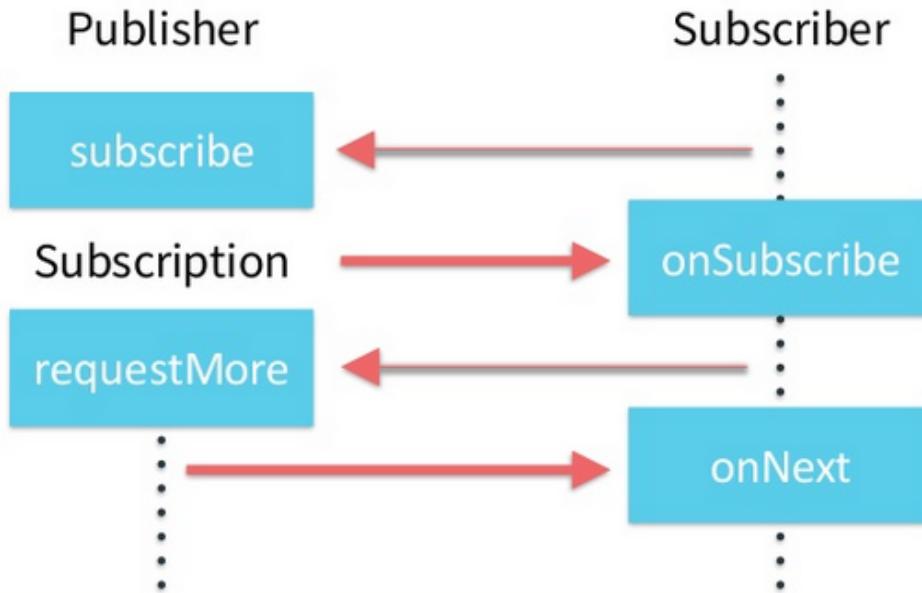
“sending 2, 5, 10, ...”

“give me max 20”

“give me max 10 more”

# How does it work?

Don't try this at home



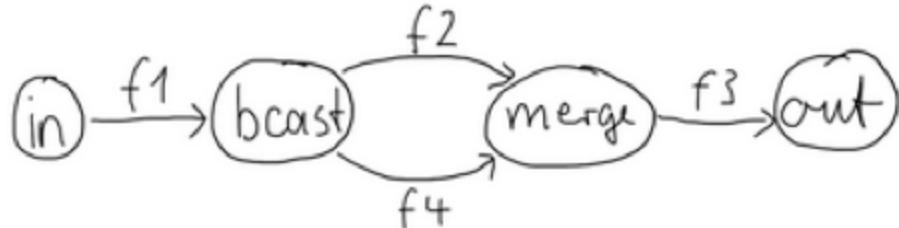
```
trait Publisher[T] {
  def subscribe(sub: Subscriber[T])
}

trait Subscription {
  def requestMore(n: Int): Unit
  def cancel(): Unit
}

trait Subscriber[T] {
  def onSubscribe(s: Subscription): Unit
  def onNext(elem: T): Unit
  def onError(thr: Throwable): Unit
  def onComplete(): Unit
}
```

# Akka Streams

- Source ~> Flow ~> Flow ~> Sink
- MaterializedFlow





# Akka Streams : advantages

- Types (stream of T)
- makes it trivially simple :-)
- Many examples online (fast and simple)
- demo of simplistic case

# simplistic Akka Streams demo

```
import akka.actor.ActorSystem
import akka.stream.ActorFlowMaterializer
import akka.stream.scaladsl._

object Words {

    implicit val system = ActorSystem("System")
    implicit val materializer = ActorFlowMaterializer()
    import system.dispatcher

    def processWords(sinkAction: (String => Unit)): Unit = {

        val text =
            """|Lorem Ipsum is simply text
            |Lorem Ipsum has been the industry standard dummy text ever since the 1500s
            |when an unknown printer took a galley of type and scrambled it to make a type
            |specimen book"""".stripMargin

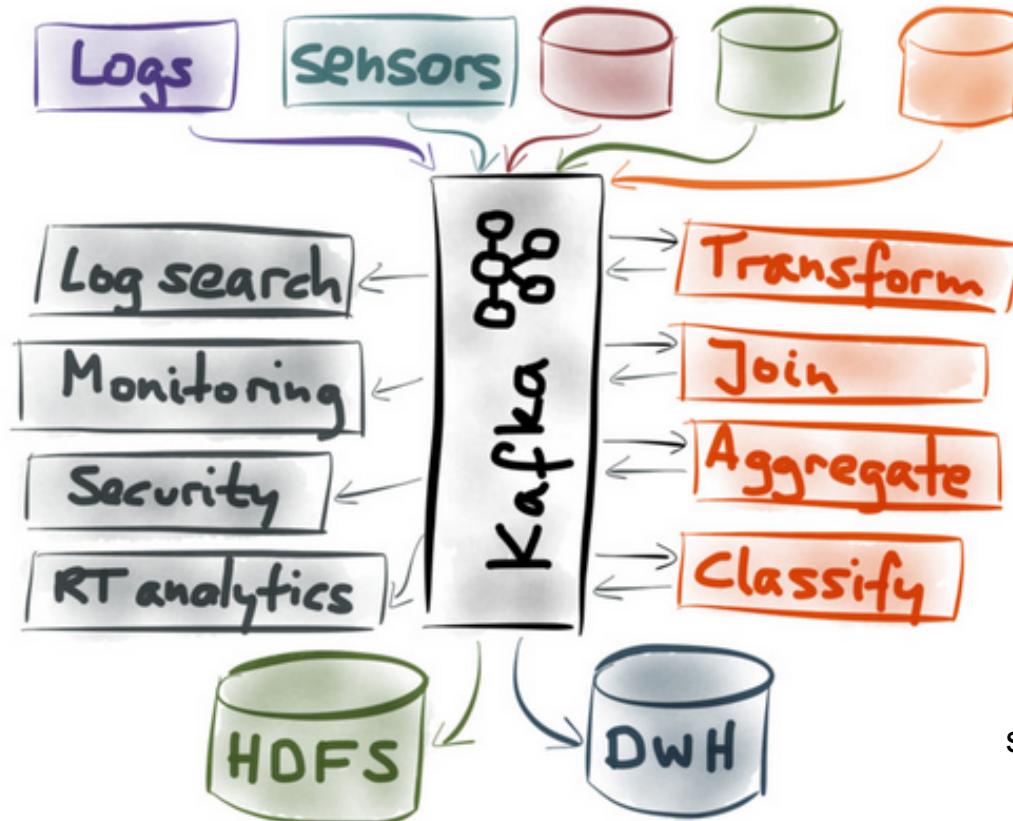
        Source(() => text.split("\s").iterator)
            .map(_.toUpperCase)
            .filter(line => line.length > 3)
            .runForeach(sinkAction)
            .onComplete(_ => system.shutdown())

        //  in ~> f1 ~> bcast ~> f2 ~> merge ~> f3 ~> out
        //          bcast ~> f4 ~> merge
    }
}
```



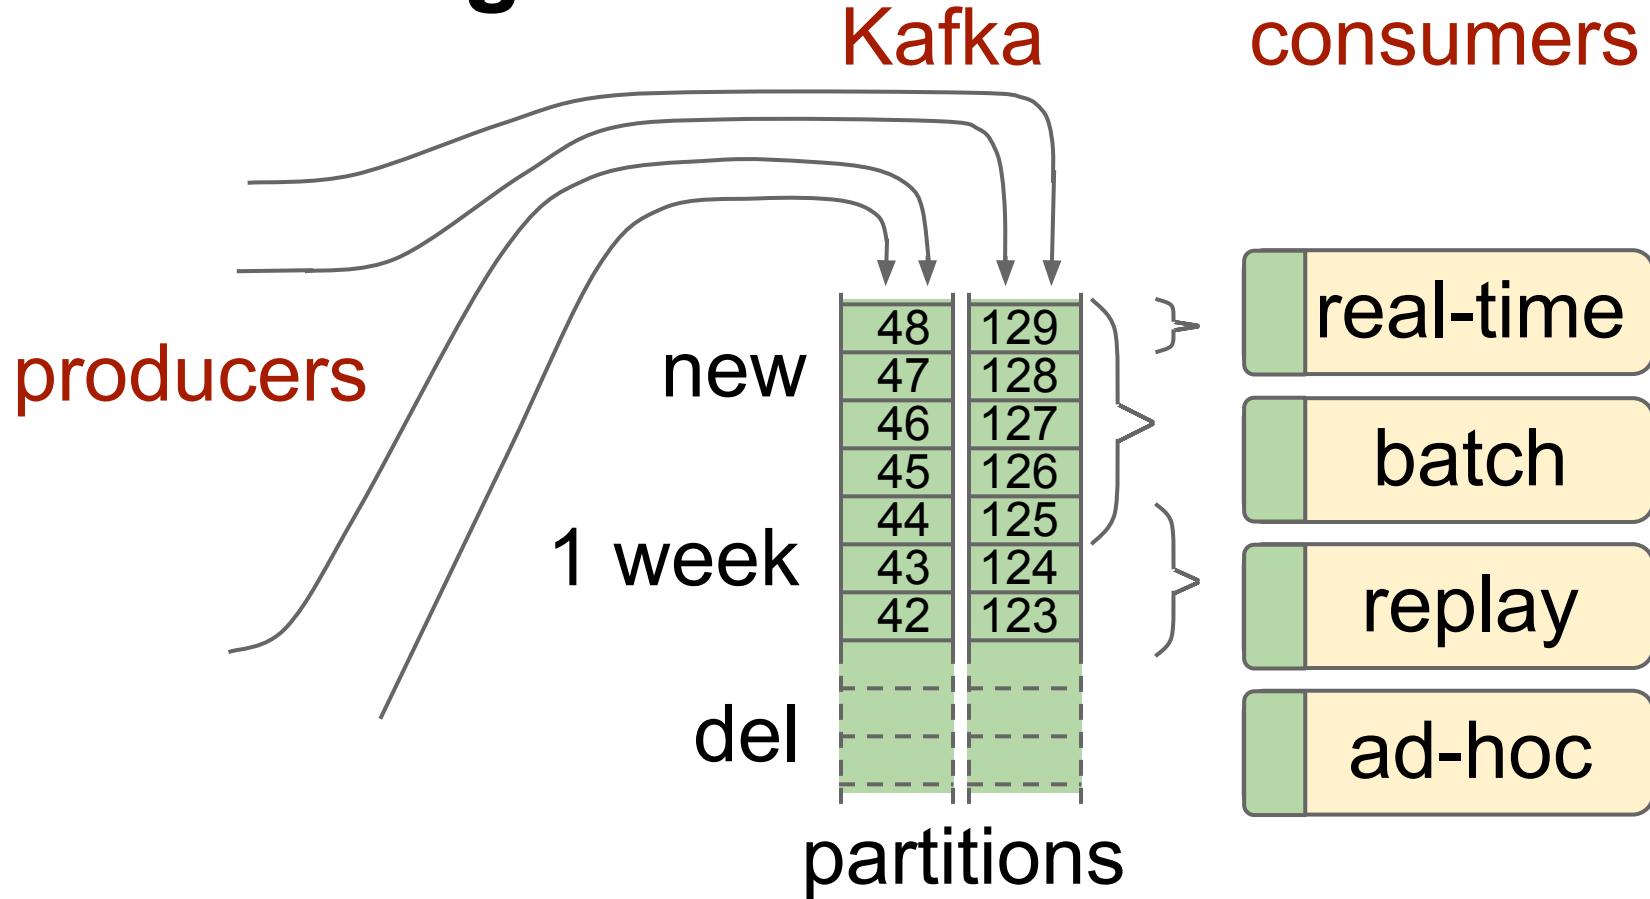
# Kafka

# Kafka (LinkedIn) : Martin Kleppmann

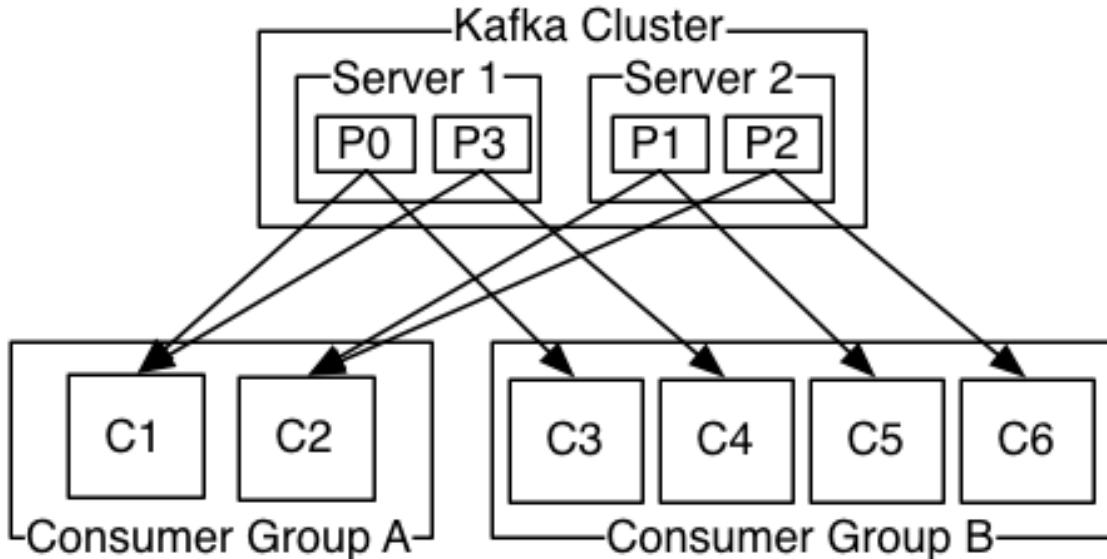


source : Martin Kleppmann  
at strata Hadoop London

# Kafka log based



# Kafka partitions



A two server Kafka cluster hosting four partitions (P0-P3) with two consumer groups. Consumer group A has two consumer instances and group B has four.



# Kafka (LinkedIn) : Jay Kreps

- 175 TB of in-flight log data per colo
  - Low-latency: ~1.5 ms
  - Replicated to each datacenter
  - Tens of thousands of data producers
  - Thousands of consumers
  - 7 million messages written/sec
  - 35 million messages read/sec
  - Hadoop integration

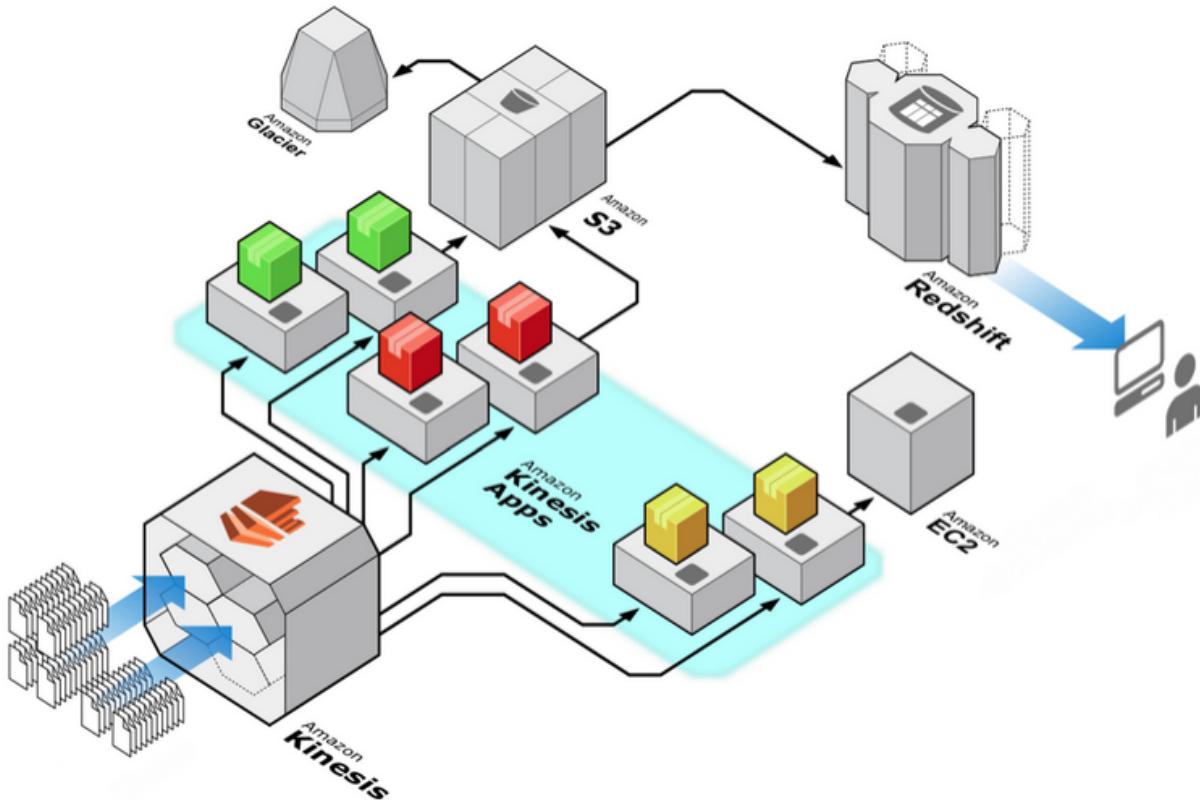
source: Jay Kreps  
on slideshare  
“I ❤️ Log”

# Real-time Data and Apache Kafka



# Kinesis

# Kinesis : Kafka as a Service



source: <http://aws.amazon.com/kinesis/details/>



# Kinesis design

- Fully (auto-)managed
- Strong durability guarantees
- Stream (= topic)
- Shard (= partition)
- “fast” writers (but ... round-trip 20 ms ?)
- “slow” readers (max 5/s per shard ??)
- Kinesis Client Library (java)



# Kinesis limitations ...

- writing latency (20 ms per entry - replicated)
  - 24 hours data retention
  - 5 reads per second

<https://brandur.org/kinesis-in-production>

- “vanishing history” after shard split
  - “if I’d understood the consequences ... earlier, I probably would have pushed harder for Kafka”

# simplistic Kinesis demo

```
import com.amazonaws.services.kinesis.AmazonKinesisClient
import com.amazonaws.services.kinesis.model.{PutRecordResult, PutRecordRequest}

class KinesisProducer(streamName: String) {

    val kinesisClient = new AmazonKinesisClient()

    def sendMessage(msg: String): String = {
        val putRequest: PutRecordRequest = new PutRecordRequest()
        putRequest.setStreamName(streamName)
        putRequest.setData(ByteBuffer.wrap(jsonMessage(msg).getBytes("UTF-8")))
        putRequest.setPartitionKey("single_partition")

        val putRecordResult: PutRecordResult = kinesisClient.putRecord(putRequest)
        putRecordResult.getSequenceNumber
    }

    // The demo Kinesis reader needs this resource + referrer format
    private def jsonMessage(referrer: String):String = {...}
}
```

Kinesis consumer with Amazon DynamoDB :: reused from <http://docs.aws.amazon.com/kinesis/latest/dev/kinesis-sample-application.html>



# Why !

(*a personal view*)

**Note: “thanks for the feedback on this section.  
Indeed Kafka and Akka serve very different purposes,  
but they both offer solutions for distributed state,  
distributed failure and slow consumers”**

# Problem 1: Distributed state



# Akka

=> state encapsulated in Actors

=> exchange self-contained messages

# Kafka

=> immutable, ordered update queue (Kappa)

# Problem 2: Distributed failure



## Akka

=> explicit failure management (supervisor)

# Kafka

=> partitions are replicated over brokers

=> consumers can replay from log

# Problem 3: Slow consumers



# Akka Streams

=> automatic back-pressure (avoid overflow)

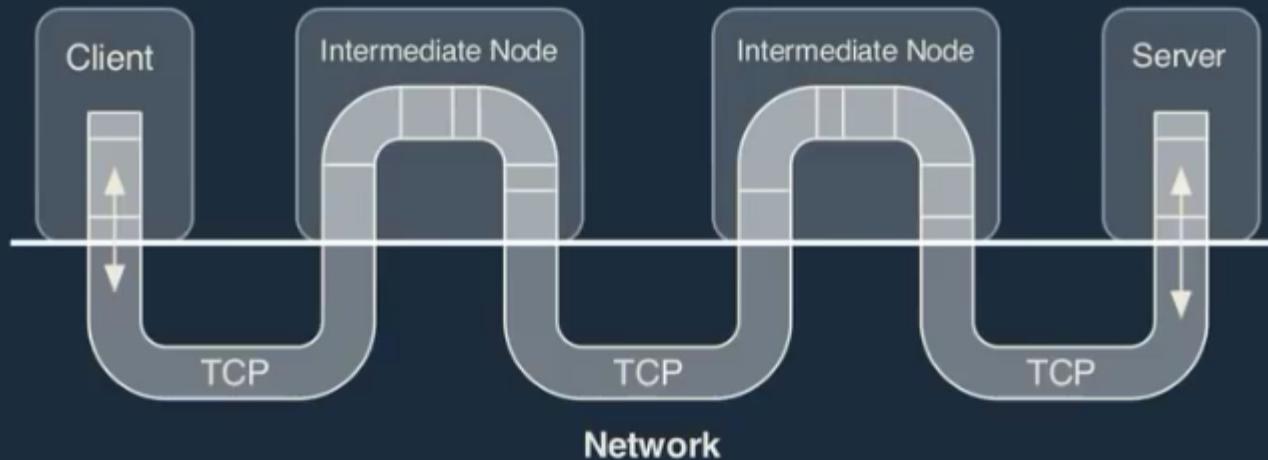
# Kafka

=> consumers fully decoupled

=> keeps data for 1 week ! (Kinesis: 1 day)

# Avoid overflow in Akka: “tight seals”

## CONTINUOUS PIPELINES ACROSS MACHINES





# Avoid overflow in Kafka: “big lake”

