



Big Data Architectural Patterns and Best Practices on AWS

Siva Raghupathy, Sr. Manager, Solutions Architecture, AWS

April 2016



Agenda

Big data challenges

How to simplify big data processing


What technologies should you use?

- Why?
- How?

Reference architecture

Design patterns

Ever Increasing Big Data



Big Data Evolution


Batch
processing

Stream
processing

Machine
learning



Plethora of Tools



logstash



IP[y]: IPython
Interactive Computing



DynamoDB
Streams

Big Data Challenges

Is there a reference architecture?

What tools should I use?

How?

Why?

Architectural Principles

Decoupled “data bus”

- Data → **Store** → **Process** → **Store** → **Analyze** → Answers

Use the right tool for the job

- Data structure, latency, throughput, access patterns

Use Lambda architecture ideas


- Immutable (append-only) log, batch/speed/serving layer

Leverage AWS managed services

- Scalable/elastic, available, reliable, secure, no/low admin

Big data ≠ big cost

Simplify Big Data Processing






COLLECT

COLLECT

STORE

Types of Data




Database records

Search documents

Log files

Messaging events

Devices / sensors / IoT stream




Store

COLLECT

STORE

Message & Stream Storage



Amazon SQS

- Managed message queue service

Apache Kafka

- High throughput distributed messaging system

Amazon Kinesis Streams

- Managed stream storage + processing

Amazon Kinesis Firehose

- Managed data delivery


Amazon DynamoDB

- Managed NoSQL database
- Tables can be stream-enabled

Why Stream Storage?


Decouple producers & consumers
Persistent buffer
Collect multiple streams

Preserve client ordering
Streaming MapReduce
Parallel consumption



What About Queues & Pub/Sub ?

- Decouple producers & consumers/subscribers
- Persistent buffer
- Collect multiple streams
- **No** client ordering
- **No** parallel consumption for Amazon SQS
 - Amazon SNS can route to multiple queues or λ functions
- **No** streaming MapReduce



What Stream Storage should I use?

	Amazon DynamoDB Streams	Amazon Kinesis Streams	Amazon Kinesis Firehose	Apache Kafka	Amazon SQS
AWS managed service	Yes	Yes	Yes	No	Yes
Guaranteed ordering	Yes	Yes	Yes	Yes	No
Delivery	exactly-once	at-least-once	exactly-once	at-least-once	at-least-once
Data retention period	24 hours	7 days	N/A	Configurable	14 days
Availability	3 AZ	3 AZ	3 AZ	Configurable	3 AZ
Scale / throughput	No limit / ~ table IOPS	No limit / ~ shards	No limit / automatic	No limit / ~ nodes	No limits / automatic
Parallel clients	Yes	Yes	No	Yes	No
Stream MapReduce	Yes	Yes	N/A	Yes	N/A
Row/object size	400 KB	1 MB	Destination row/object size	Configurable	256 KB
Cost	Higher (table cost)	Low	Low	Low (+admin)	Low-medium

Hot

Warm

COLLECT

STORE

File Storage

Applications

Web apps 

Mobile apps 

Data centers AWS Direct Connect 

Logging

Logging 

AWS CloudTrail  Amazon CloudWatch 

Transport

AWS Import/Export  Snowball


Messaging

Message 

IoT

Devices 

AWS IoT 



Amazon S3


Why is Amazon S3 Good for Big Data?

- Natively supported by big data frameworks (Spark, Hive, Presto, etc.)
- No need to run compute clusters for storage (unlike HDFS)
- Can run transient Hadoop clusters & Amazon EC2 Spot Instances
- Multiple distinct (Spark, Hive, Presto) clusters can use the same data
- Unlimited number of objects
- Very high bandwidth – no aggregate throughput limit
- Highly available – can tolerate AZ failure
- Designed for 99.99999999% durability
- Tired-storage (Standard, IA, Amazon Glacier) via life-cycle policy
- Secure – SSL, client/server-side encryption at rest
- Low cost

What about HDFS & Amazon Glacier?


- Use HDFS for very frequently accessed (hot) data
- Use Amazon S3 Standard for frequently accessed data
- Use Amazon S3 Standard – IA for infrequently accessed data
- Use Amazon Glacier for archiving cold data






Cache, database, search


Database Anti-pattern



Best Practice - Use the Right Tool for the Job



Materialized Views



What Data Store Should I Use?

Data structure → Fixed schema, JSON, key-value

Access patterns → Store data in the format you will access it

Data / access characteristics → Hot, warm, cold


Cost → Right cost

Data Structure and Access Patterns

Access Patterns	What to use?
Put/Get (key, value)	Cache, NoSQL
Simple relationships → 1:N, M:N	NoSQL
Cross table joins, transaction, SQL	SQL
Faceting, search	Search

Data Structure	What to use?
Fixed schema	SQL, NoSQL
Schema-free (JSON)	NoSQL, Search
(Key, value)	Cache, NoSQL


What Is the Temperature of Your Data / Access ?



Data / Access Characteristics: Hot, Warm, Cold

	Hot	Warm	Cold
Volume	MB–GB	GB–TB	PB
Item size	B–KB	KB–MB	KB–TB
Latency	ms	ms, sec	min, hrs
Durability	Low – high	High	Very high
Request rate	Very high	High	Low
Cost/GB	\$\$-\$	\$-¢¢	¢





What Data Store Should I Use?

	Amazon ElastiCache	Amazon DynamoDB	Amazon RDS/Aurora	Amazon Elasticsearch	Amazon S3	Amazon Glacier
Average latency	ms	ms	ms, sec	ms,sec	ms,sec,min (~ size)	hrs
Typical data stored	GB	GB–TBs (no limit)	GB–TB (64 TB max)	GB–TB	MB–PB (no limit)	GB–PB (no limit)
Typical item size	B-KB	KB (400 KB max)	KB (64 KB max)	KB (2 GB max)	KB-TB (5 TB max)	GB (40 TB max)
Request Rate	High – very high	Very high (no limit)	High	High	Low – high (no limit)	Very low
Storage cost GB/month	\$\$	¢¢	¢¢	¢¢	¢	¢/10
Durability	Low - moderate	Very high	Very high	High	Very high	Very high
Availability	High 2 AZ	Very high 3 AZ	Very high 3 AZ	High 2 AZ	Very high 3 AZ	Very high 3 AZ
Hot data			Warm data			Cold data

Cost Conscious Design

Example: Should I use Amazon S3 or Amazon DynamoDB?

“I’m currently scoping out a project that will greatly increase my team’s use of Amazon S3. Hoping you could answer some questions. The current iteration of the design calls for **many small files**, perhaps up to a **billion during peak**. The **total size** would be on the order of **1.5 TB per month...**”

Request rate (Writes/sec)	Object size (Bytes)	Total size (GB/month)	Objects per month
300	2048	1483	777,600,000

Cost Conscious Design

Example: Should I use Amazon S3 or Amazon DynamoDB?



Simple Monthly
Calculator

<https://calculator.s3.amazonaws.com/index.html>

Amazon S3 or Amazon DynamoDB?

Request rate (Writes/sec)	Object size (Bytes)	Total size (GB/month)	Objects per month
300	2,048	1,483	777,600,000

Amazon DynamoDB is a high performance non-relational database service that is easy to set up, operate, and scale. It is designed to address the core problems of database management, performance, scalability, and reliability. It also provides predictable high performance and low latency at scale.

Indexed Data Storage:

Dataset Size:

1483 GB

Provisioned Throughput Capacity *:

Item Size (All attributes):

2 KB

Number of items read per second:

0 Reads/Second

Read Consistency:

Strongly Consistent

Eventually Consistent (cheaper)

Number of items written per second:

300 Writes/Second

Amazon S3 is storage for the Internet. It is designed to make web-scale computing easier for developers.

Storage:

Storage:

1483 GB

Reduced Redundancy Storage:

0 GB

Requests:

PUT/COPY/POST/LIST Requests:

77760000 Requests

GET and Other Requests:

0 Requests

Amazon DynamoDB Service (US-East)

\$ 644.30

Provisioned Throughput Capacity:

\$ 261.69

Indexed Data Storage:

\$ 382.61

Amazon S3 Service (US-East)

\$ 3932.27

Storage:

\$ 44.27

Put/List Requests:

\$ 3888.00




SIMPLE MONTHLY CALCULATOR



use

<u>Amazon S3 Service (US-East)</u>		\$ 3932.27
Storage:	\$ 44.27	
Put/List Requests:	\$ 3888.00	
<u>Amazon DynamoDB Service (US-East)</u>		\$ 644.30
Provisioned Throughput Capacity:	\$ 261.69	
Indexed Data Storage:	\$ 382.61	
DynamoDB Streams:	\$ 0.00	

	Request rate (Writes/sec)	Object size (Bytes)	Total size (GB/month)	Objects per month
<u>Scenario 1</u>	300	2,048	1,483	777,600,000
<u>Scenario 2</u>	300	32,768	23,730	777,600,000



Amazon S3

use

<u>Amazon S3 Service (US-East)</u>		\$ 4588.55
Storage:	\$ 700.55	
Put/List Requests:	\$ 3888.00	
<u>Amazon DynamoDB Service (US-East)</u>		\$ 10131.40
Provisioned Throughput Capacity:	\$ 4187.04	
Indexed Data Storage:	\$ 5944.36	
DynamoDB Streams:	\$ 0.00	



Process /
analyze



PROCESS /
ANALYZE

Process / Analyze

- Batch - Minutes or hours on cold data
 - Daily/weekly/monthly reports
- Interactive – Seconds on warm/cold data
 - Self-service dashboards
- Messaging – Milliseconds or seconds on hot data
 - Message/event buffering
- Streaming - Milliseconds or seconds on hot data
 - Billing/fraud alerts, 1 minute metrics

Predictions via Machine Learning

ML gives computers the ability to learn without being explicitly programmed

Machine learning algorithms:

Supervised learning ← “teach” program

- Classification ← Is this transaction fraud? (yes/no)
- Regression ← Customer life-time value?

Unsupervised learning ← Let it learn by itself

- Clustering ← Market segmentation

Tools and Frameworks

Machine Learning

- Amazon ML, Amazon EMR (Spark ML)

Interactive

- Amazon Redshift, Amazon EMR (Presto, Spark)

Batch


- Amazon EMR (MapReduce, Hive, Pig, Spark)

Messaging

- Amazon SQS application on Amazon EC2

Streaming

- Micro-batch: Spark Streaming, KCL
- Real-time: Amazon Kinesis Analytics, Storm, AWS Lambda, KCL



What Streaming / Messaging Technology Should I Use?

	Spark Streaming	Apache Storm	Kinesis KCL Application	AWS Lambda	Amazon SQS Apps
Scale	~ Nodes	~ Nodes	~ Nodes	Automatic	~ Nodes
Micro-batch or Real-time	Micro-batch	Real-time	Near-real-time	Near-real-time	Near-real-time
AWS managed service	Yes (EMR)	No (EC2)	No (KCL + EC2 + Auto Scaling)	Yes	No (EC2 + Auto Scaling)
Scalability	No limits ~ nodes	No limits ~ nodes	No limits ~ nodes	No limits	No limits
Availability	Single AZ	Configurable	Multi-AZ	Multi-AZ	Multi-AZ
Programming languages	Java, Python, Scala	Any language via Thrift	Java, via MultiLang Daemon (.NET, Python, Ruby, Node.js)	Node.js, Java, Python	AWS SDK languages (Java, .NET, Python, ...)
	Fast	Fast	Fast	Fast	Fast

What Analytics Technology Should I Use?

	Amazon Redshift	Amazon EMR		
		Presto	Spark	Hive
Query latency	Low	Low	Low	High
Durability	High	High	High	High
Data volume	1.6 PB max	~Nodes	~Nodes	~Nodes
AWS managed	Yes	Yes	Yes	Yes
Storage	Native	HDFS / S3	HDFS / S3	HDFS / S3
SQL compatibility	High	High	Low (SparkSQL)	Medium (HQL)
	Fast	Fast	Fast	Slow

What About ETL?



Data Integration

Reduce the effort to move, cleanse, synchronize, manage, and automatize data related processes.

ATTUNITY

Attunity CloudBeam

informatica

Informatica Cloud

Matillion
Business Intelligence


Matillion ETL for Redshift

snapLogic

snapLogic

alteryx

alteryx





CONSUME




Applications & API


Analysis and visualization

Notebooks

IDE



Putting It All Together



Design Patterns


Primitive: Multi-Stage Decoupled “Data Bus”

Multiple stages

Storage decouples multiple processing stages




Primitive: Multiple Stream Processing Applications Can Read from Amazon Kinesis




process
 store

Primitive: Analysis Frameworks Could Read from Multiple Data Stores




process
store


Data Temperature vs. Processing Speed




Real-time Analytics




Message / Event Processing



Interactive & Batch Analytics



Lambda Architecture



Summary

Decoupled “data bus”

- Data → **Store** → **Process** → **Store** → **Analyze** → Answers

Use the right tool for the job

- Data structure, latency, throughput, access patterns

Use Lambda architecture ideas

- Immutable (append-only) log, batch/speed/serving layer

Leverage AWS managed services

- Scalable/elastic, available, reliable, secure, no/low admin

Big data ≠ big cost



Thank you!

aws.amazon.com/big-data